

Homework 6

PSTAT 131/231

Contents

```
library(tidyverse)
library(tidymodels)
library(rpart.plot)
library(ISLR)
library(vip)
library(randomForest)
library(xgboost)
```

Exercise 1

Read in the data and set things up as in Homework 5:

- Use `clean_names()`
- Filter out the rarer Pokémon types
- Convert `type_1` and `legendary` to factors

```
library(janitor)
pokemon <- read.csv('Pokemon.csv')
pokemon <- janitor:: clean_names(dat = pokemon)
pokemon1 <- pokemon %>%
  filter(type_1 %in% c('Bug','Fire', 'Grass', 'Normal', 'Water', 'Psychic'))  
  
pokemon2 <- pokemon1 %>%
  mutate(type_1 = factor(type_1), legendary = factor(legendary),
         generation = factor(generation))
```

Do an initial split of the data; you can choose the percentage for splitting. Stratify on the outcome variable.

```
set.seed(3435)
pokemon_split <- initial_split(pokemon2, prop = 0.7,
                                 strata = type_1)
pokemon_train <- training(pokemon_split)
pokemon_test <- testing(pokemon_split)
```

Fold the training set using v -fold cross-validation, with $v = 5$. Stratify on the outcome variable.

```
pokemon_folds <- vfold_cv(pokemon_train, v = 5, strata = type_1)
```

Set up a recipe to predict `type_1` with `legendary`, `generation`, `sp_atk`, `attack`, `speed`, `defense`, `hp`, and `sp_def`:

- Dummy-code `legendary` and `generation`;
- Center and scale all predictors.


```
pokemon_recipe <-  
  recipe(type_1 ~ legendary + generation + sp_atk + attack + speed + defense + hp + sp_def,  
         data = pokemon_train) %>%  
  step_dummy(legendary, generation) %>%  
  step_center(all_predictors()) %>%  
  step_scale(all_predictors())
```

Exercise 2

Create a correlation matrix of the training set, using the `corrplot` package. *Note: You can choose how to handle the continuous variables for this plot; justify your decision(s).*

```
library(corr)
cor_pokemon <- pokemon_train %>%
  select(-x) %>%
  select(is.numeric) %>%
  correlate()

cor_pokemon %>%
  stretch() %>%
  ggplot(aes(x,y,fill = r)) +
  geom_tile() +
  geom_text(aes(label = as.character(fashion(r))))
```



What relationships, if any, do you notice? Do these relationships make sense to you?

I choose to remove the variable `x`, because it is just the ID of each Pokemon. Also, I remove other categorical value. In the plot above, we could notice that total is positive correlate to other variables except for generation. These relationships make sense to me. Total is a general guide to how strong a Pokemon is, so each if each attributes is high and the total would be high.

Exercise 3

First, set up a decision tree model and workflow. Tune the `cost_complexity` hyperparameter. Use the same levels we used in Lab 7 – that is, `range = c(-3, -1)`. Specify that the metric we want to optimize is `roc_auc`.

```
tree_spec <- decision_tree() %>%
  set_engine("rpart") %>%
  set_mode("classification") %>%
  set_args(cost_complexity = tune())

class_tree_wf <- workflow() %>%
  add_model(tree_spec) %>%
  add_recipe(pokemon_recipe)

param_grid <- grid_regular(cost_complexity(range = c(-3, -1)), levels = 10)

tune_res <- tune_grid(
  class_tree_wf,
```


```

    resamples = pokemon_folds,
    grid = param_grid,
    metrics = metric_set(roc_auc)
)

```

Print an `autoplot()` of the results. What do you observe? Does a single decision tree perform better with a smaller or larger complexity penalty?

```
autoplot(tune_res)
```



I notice when cost-complexity parameter increase, our `roc_auc` increase. However, when cost-complexity parameter close to 0.07, `roc_auc` reaches to peak and decrease sharply. A single decision tree perform better with a smaller complexity penalty.

Exercise 4

What is the `roc_auc` of your best-performing pruned decision tree on the folds? Hint: Use `collect_metrics()` and `arrange()`.

```

tree_roc_auc <- collect_metrics(tune_res) %>%
  arrange(-mean)
tree_roc_auc <- tree_roc_auc %>% head(1)
tree_roc_auc

```

```
## # A tibble: 1 x 7
```

```

##   cost_complexity .metric .estimator  mean      n std_err .config
##                 <dbl> <chr>    <chr>     <dbl> <int>   <dbl> <chr>
## 1          0.0215 roc_auc hand_till  0.629      5  0.0191 Preprocessor1_Model07

```

The roc_auc of best-performing pruned decision tree on the folds is 0.628788

Exercise 5

Using `rpart.plot`, fit and visualize your best-performing pruned decision tree with the *training* set.

```


best_complexity <- select_best(tune_res)

class_tree_final <- finalize_workflow(class_tree_wf, best_complexity)

class_tree_final_fit <- fit(class_tree_final, data = pokemon_train)

class_tree_final_fit %>%
  extract_fit_engine() %>%
  rpart.plot()

```



Exercise 5

Now set up a random forest model and workflow. Use the `ranger` engine and set `importance = "impurity"`. Tune `mtry`, `trees`, and `min_n`. Using the documentation for `rand_forest()`, explain in your own words

what each of these hyperparameters represent.

```
rf_spec <- rand_forest(mtry = tune(), trees = tune(), min_n = tune()) %>%
  set_engine("ranger", importance = "impurity") %>%
  set_mode("classification")

rf_wf <- workflow() %>%
  add_recipe(pokemon_recipe) %>%
  add_model(rf_spec)
```

- mtry means when creating a tree model how many predictors will be randomly sampled at each split
- trees means the number of trees contained in the ensemble
- min_n means the minimum number of data points in a node that are required for the node to be split further.

Create a regular grid with 8 levels each. You can choose plausible ranges for each hyperparameter. Note that `mtry` should not be smaller than 1 or larger than 8. **Explain why not. What type of model would `mtry = 8` represent?**

```
rp_grid <- grid_regular(mtry(range = c(2,7)), trees(range = c(10,2000)),
                         min_n(range = c(2, 10)), levels = 8)
```


`mtry` should not be smaller than 1 or larger than 8 because we need use at least one predictor in our model and total number of predictors are 8 so `mtry` could not be larger than 8. When `mtry = 8`, we use all the predictors and the type of model is bagging.

Exercise 6

Specify `roc_auc` as a metric. Tune the model and print an `autoplot()` of the results. What do you observe? What values of the hyperparameters seem to yield the best performance?

```
tune_rf <- tune_grid(
  rf_wf,
  resamples = pokemon_folds,
  grid = rp_grid,
  metrics = metric_set(roc_auc)
)

autoplot(tune_rf)
```



I observed that when number of trees increases to 250, our roc_auc increase sharply, and then when number of trees increases, our roc_auc does not have a distinct increase. The number of predictors choice has no large influence on our roc_auc, the overall trend seems similar to each other. The choice of minimal node size seems have no large effect on our model accuracy.

Exercise 7

What is the roc_auc of your best-performing random forest model on the folds? Hint: Use `collect_metrics()` and `arrange()`.

```
rf_roc_auc <- collect_metrics(tune_rf) %>%
  arrange(-mean)
rf_roc_auc <- rf_roc_auc %>% head(1)
rf_roc_auc

## # A tibble: 1 x 9
##   mtry trees min_n .metric .estimator mean    n std_err .config
##   <int> <int> <int> <chr>   <chr>    <dbl> <int> <dbl> <chr>
## 1     4    1147     3 roc_auc hand_till  0.729      5  0.0139 Preprocessor1_Model0~
```

The roc_auc of my best-performing model on the folds is 0.729108


Exercise 8

Create a variable importance plot, using `vip()`, with your best-performing random forest model fit on the *training* set.

Which variables were most useful? Which were least useful? Are these results what you expected, or not?

```
best_rf <- select_best(tune_rf, metric = "roc_auc")
rf_final <- finalize_workflow(rf_wf, best_rf)
rf_final_fit <- fit(rf_final, data = pokemon_train)

rf_final_fit %>%
  extract_fit_engine() %>%
  vip()
```



Sp_atk and attack were most useful, and generation and legendary were least useful. Some other attributes like hp speed defense were still important. These results are what I expected because type could determine Pokemon resistant to others.

Exercise 9

Finally, set up a boosted tree model and workflow. Use the `xgboost` engine. Tune `trees`. Create a regular grid with 10 levels; let `trees` range from 10 to 2000. Specify `roc_auc` and again print an `autoplot()` of the results.

```
boost_spec <- boost_tree(trees = tune()) %>%
  set_engine("xgboost") %>%
  set_mode("classification")

boost_wf <- workflow() %>%
```

```


add_recipe(pokemon_recipe) %>%
  add_model(boost_spec)

trees_grid2 <- grid_regular(trees(range = c(10,2000)), levels = 10)

tune_bt <- tune_grid(
  boost_wf,
  resamples = pokemon_folds,
  grid = trees_grid2,
  metrics = metric_set(roc_auc)
)

autoplot(tune_bt)

```



What do you observe?

The roc_auc reach to the peak at #Trees around 250, and then has a decreasing trend.

What is the roc_auc of your best-performing boosted tree model on the folds? *Hint: Use collect_metrics() and arrange().*

```

boost_roc_auc <- collect_metrics(tune_bt) %>%
  arrange(-mean)
boost_roc_auc <- boost_roc_auc %>% head(1)
boost_roc_auc

```

```

## # A tibble: 1 x 7
##   trees .metric .estimator  mean     n std_err .config
##   <int> <chr>    <chr>     <dbl> <int>   <dbl> <chr>
## 1    231 roc_auc hand_till  0.698     5  0.0139 Preprocessor1_Model02

```

The roc_auc of my best-performing boosted tree model on the fold is 0.6977248

Exercise 10

Display a table of the three ROC AUC values for your best-performing pruned tree, random forest, and boosted tree models. Which performed best on the folds? Select the best of the three and use `select_best()`, `finalize_workflow()`, and `fit()` to fit it to the *testing* set.

```

model_roc_auc <- c(tree_roc_auc$mean, rf_roc_auc$mean, boost_roc_auc$mean)
model <- c("pruned tree", "random forest", "boosted tree")
result <- tibble(model_roc_auc, model)
result %>% arrange(-model_roc_auc)

```

```

## # A tibble: 3 x 2
##   model_roc_auc model
##       <dbl> <chr>
## 1      0.729 random forest
## 2      0.698 boosted tree
## 3      0.629 pruned tree

```

- random forest performed best on the folds

```

rf_parameters <- select_best(tune_rf, metric = "roc_auc")
rf_final_test <- finalize_workflow(rf_wf, rf_parameters)
rf_final_testfit <- fit(rf_final_test, data = pokemon_test)

```


Print the AUC value of your best-performing model on the testing set. Print the ROC curves. Finally, create and visualize a confusion matrix heat map.

```

predictions_result <- augment(rf_final_fit, new_data = pokemon_test) %>%
  select(type_1, .pred_class, .pred_Bug, .pred_Fire, .pred_Grass, .pred_Normal, .pred_Psychic,
         .pred_Water)

predictions_result %>%
  roc_curve(type_1, .pred_Bug, .pred_Fire, .pred_Grass, .pred_Normal, .pred_Psychic,
            .pred_Water) %>%
  autoplot()

```



```
predictions_result %>% conf_mat(type_1, .pred_class) %>%
  autoplot(type = "heatmap")
```

	Bug	Fire	Grass	Normal	Psychic	Water
Bug -	9	0	2	2	0	3
Fire -	0	4	2	1	1	6
Grass -	0	2	1	0	1	0
Normal -	5	2	1	20	1	7
Psychic -	1	3	5	1	10	5
Water -	6	5	10	6	5	13
Truth	Bug	Fire	Grass	Normal	Psychic	Water

Which classes was your model most accurate at predicting? Which was it worst at?

My model was most accurate at predicting Normal, and it was worst at predicting class Water.